

Efficient encodings for inverted index

Yiliang Xiong

April 19, 2018

Overview

Variable-byte encoding

Simple9/Simple16 encoding

Frame-Of-Reference encoding

Elias-Fano encoding

Overview

Posting list

A posting list can be defined as a monotonically increasing sequence

$$0 \leq x_0 \leq x_1, \dots, x_{n-1} \leq u$$

where u is the upper bound of *doc id*.

Posting list

A posting list can be defined as a monotonically increasing sequence

$$0 \leq x_0 \leq x_1, \dots, x_{n-1} \leq u$$

where u is the upper bound of *doc id*.

The sequence can also be represented by the gaps between consecutive numbers:

$$\delta_0, \delta_1, \dots, \delta_{n-1}$$

where $\delta_0 = 0$ and $\delta_i = x_i - x_{i-1}$.

Posting list

A posting list can be defined as a monotonically increasing sequence

$$0 \leq x_0 \leq x_1, \dots, x_{n-1} \leq u$$

where u is the upper bound of *doc id*.

The sequence can also be represented by the gaps between consecutive numbers:

$$\delta_0, \delta_1, \dots, \delta_{n-1}$$

where $\delta_0 = 0$ and $\delta_i = x_i - x_{i-1}$.

How to restore the original sequence?

Posting list

A posting list can be defined as a monotonically increasing sequence

$$0 \leq x_0 \leq x_1, \dots, x_{n-1} \leq u$$

where u is the upper bound of *doc id*.

The sequence can also be represented by the gaps between consecutive numbers:

$$\delta_0, \delta_1, \dots, \delta_{n-1}$$

where $\delta_0 = 0$ and $\delta_i = x_i - x_{i-1}$.

How to restore the original sequence?


- ▶ Prefix sum $x_i = \sum_0^i \delta_i$
- ▶ High-performance algorithm exists:
https://en.wikipedia.org/wiki/Prefix_sum#Parallel_algorithm

Variable-byte encoding


Google's index format in early days

Byte-Aligned Variable-length Encodings

- Varint encoding:
 - 7 bits per byte with continuation bit
 - Con: Decoding requires lots of branches/shifts/masks



- Idea: Encode byte length as low 2 bits
 - Better: fewer branches, shifts, and masks
 - Con: Limited to 30-bit values, still some shifting to decode




Stolen from Jeff Dean's Challenges in Building Large-Scale Information Retrieval Systems.

Google's index format in early days (cont'd)

Group Varint Encoding

- Idea: encode groups of 4 values in 5-17 bytes
 - Pull out 4 2-bit binary lengths into single byte prefix



- Decode: Load prefix byte and use value to lookup in 256-entry table:
...
`00000110` → Offsets: +1,+2,+3,+5; Masks: ff, ff, ffff, ffffff
...
...
- Much faster than alternatives:
 - 7-bit-per-byte varint: decode ~180M numbers/second
 - 30-bit Varint w/ 2-bit length: decode ~240M numbers/second
 - Group varint: decode ~400M numbers/second

Stolen from Jeff Dean's Challenges in Building Large-Scale Information Retrieval Systems.

More examples: Protobuf encoding

Each byte in a varint, except the last byte, has the most significant bit (msb) set. This indicates that there are further bytes to come.

So, for example, here is the number 1 - it's a single byte, so the msb is not set:

00000001

And here is 300 - this is a bit more complicated:

10101100 00000010

More details can be found at <https://developers.google.com/protocol-buffers/docs/encoding>

More examples: UTF-8

Bits at the beginning of the first byte determine the code pattern:

- ▶ 0_____ backwards compatibility with ASCII
- ▶ 11_____ the first byte of a UTF-8 encoded character
- ▶ 10_____ the remaining byte of a UTF-8 encoded character

Number of bytes	Bits for code point	First code point	Last code point	Byte 1	Byte 2	Byte 3	Byte 4
1	7	U+0000	U+007F	0xxxxxxx			
2	11	U+0080	U+07FF	110xxxxx	10xxxxxx		
3	16	U+0800	U+FFFF	1110xxxx	10xxxxxx	10xxxxxx	
4	21	U+10000	U+10FFFF	11110xxx	10xxxxxx	10xxxxxx	10xxxxxx

See history of UTF-8 at <https://www.cl.cam.ac.uk/~mgk25/ucs/utf-8-history.txt>

Simple9/Simple16 encoding

Simple9 (S9) encoding

The basic idea is to try to pack as many integers as possible into one 32-bit word.

To do this, Simple9 divides each word into 4 status bits and 28 data bits, where the data bits can be divided up in 9 different ways:

- 28 1-bit numbers
- 14 2-bit numbers
- 9 3-bit numbers (onebit unused)
- 7 4-bit numbers
- 5 5-bit numbers (three bits unused)
- 4 7-bit numbers
- 3 9-bit numbers (one bit unused)
- 2 14-bit numbers
- 1 28-bit number

Vo Ngoc Anh, Alistair Moffat, Index Compression using Fixed Binary Codewords (ADC 2004).

Simple16 (S16) encoding

Simple9 wastes bits in two ways,

- ▶ by having only 9 cases instead of the 16 that can be expressed with 4 status bits, and
- ▶ by having unused bits in several of these cases.

Simple16 (S16) encoding

Simple9 wastes bits in two ways,

- ▶ by having only 9 cases instead of the 16 that can be expressed with 4 status bits, and
- ▶ by having unused bits in several of these cases.

A new variation that avoids this:

```
5 5-bit numbers in Simple9, with 3 bits unused ==>
3 6-bit numbers followed by 2 5-bit numbers
2 5-bit numbers followed by 3 6-bit numbers
```

Jiangong Zhang, Xiaohui Long, Torsten Suel, Performance of Compressed Inverted List Caching in Search Engines (WWW 2008)

Frame-Of-Reference encoding

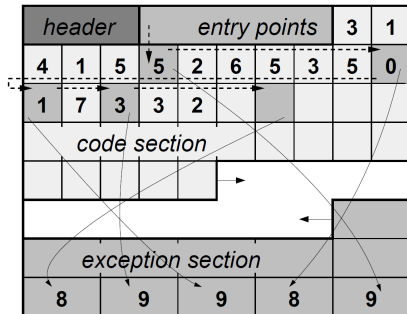
A naive approach to decompression

The naive way to implement any decompression scheme that distinguishes between coded and *exception* values, is to use a special code MAXCODE for exceptions:

```
/* NAIVE approach to decompression */
for(i = j = 0; i < n; i++) {
    if (code[i] < MAXCODE) {
        output[i] = DECODE(code[i]);
    } else {
        output[i] = exception[--j]);
    }
}
```

Patched Frame-Of-Reference (PForDelta)

An elegant and efficient way to handle exceptions:



- What if two consecutive exceptions have a distance of more than 2^b ?

Marcin Zukowski, Sandor Heman, Niels Nes, Peter Boncz, Super-Scalar RAM-CPU Cache Compression (ICDE '06).

PForDelta decompression

Compared to the naive implementation:


- ▶ no if-else statement
- ▶ reduce branch misprediction (better CPU pipelining)

```
int Decompress<ANY>(
    int n, int b, ANY *output, void *input, ANY *exception, int *next_exception) {
    int next, code[n], cur = *next_exception;
    UNPACK[b](code, input, n); /* bit-unpack the values */
    for(int i = 0; i < n; i++) { /* LOOP1: decode regardless */
        output[i] = DECODE(code[i]);
    }
    for(int i = 1; cur < n; i++, cur = next) { /* LOOP2: patch it up */
        next = cur + code[cur] + 1;
        output[cur] = exception[-i];
    }
    *next_exception = cur - n;
    return i;
}
```

Elias-Fano encoding

Elias-Fano encoding

- ▶ Given n and u we have a monotonically increasing sequence $0 \leq x_0 \leq x_1, \dots, x_{n-1} \leq u$
- ▶ Store the lower $l = \log(u/n)$ bits explicitly
- ▶ Store the upper bits as a sequence of unary coded gaps ($0^k 1$ represents k)
- ▶ Use at most $2 + \log(u/n)$ bits per element



Peter Elias, Efficient Storage and Retrieval by Content and Address of Static Files (JACM 1974).

Sebastiano Vigna, Quasi-Succinct Indices (WSDM '13).

How to search

Looking up k -th element:

- ▶ Suppose you want to get the k -th element quickly
- ▶ Just scan the upper bits, one word at a time, doing population counting
- ▶ When you get to the right word, complete sequentially and pick the lower bits

How to search

Looking up k-th element:

- ▶ Suppose you want to get the k-th element quickly
- ▶ Just scan the upper bits, one word at a time, doing population counting
- ▶ When you get to the right word, complete sequentially and pick the lower bits

Searching for any element $x_i \geq q$:

- ▶ It's exactly the same: only, you count zeroes
- ▶ Zeroes tells you how much the upper bits are increasing, which is the important thing
- ▶ Just skip $q \gg 1$ upper zeroes and complete sequentially

Advantages

Almost optimal space usage

- ▶ $2 + \log(u/n)$ bits per element

Advantages

Almost optimal space usage

- ▶ $2 + \log(u/n)$ bits per element

Independent of *doc id* distribution

- ▶ works for long and short posting list

Advantages

Almost optimal space usage

- ▶ $2 + \log(u/n)$ bits per element

Independent of *doc id* distribution

- ▶ works for long and short posting list

Reading sequentially requires very few logical operations

- ▶ only shift and bitwise operation

Advantages

Almost optimal space usage

- ▶ $2 + \log(u/n)$ bits per element

Independent of *doc id* distribution

- ▶ works for long and short posting list

Reading sequentially requires very few logical operations

- ▶ only shift and bitwise operation

Encoding integer sequence directly instead of $\delta_i = x_i - x_{i-1}$

- ▶ no need to calculate *prefix sum* after decoding

Thank you

Any questions?